

solution2

February 15, 2020

1 Machine, Data and Learning: Assignment 1

Assignment 1: Question 2

Animesh Sinha (2018113001)

Mukund Choudhary (2018114015)

1.1 Loading the Data

We start by loading the data into Numpy Arrays and Pandas DataFrames, and plotting the basic structure of it all. Some sampling, shuffling and splitting is done as we were tasked in the problem statement.

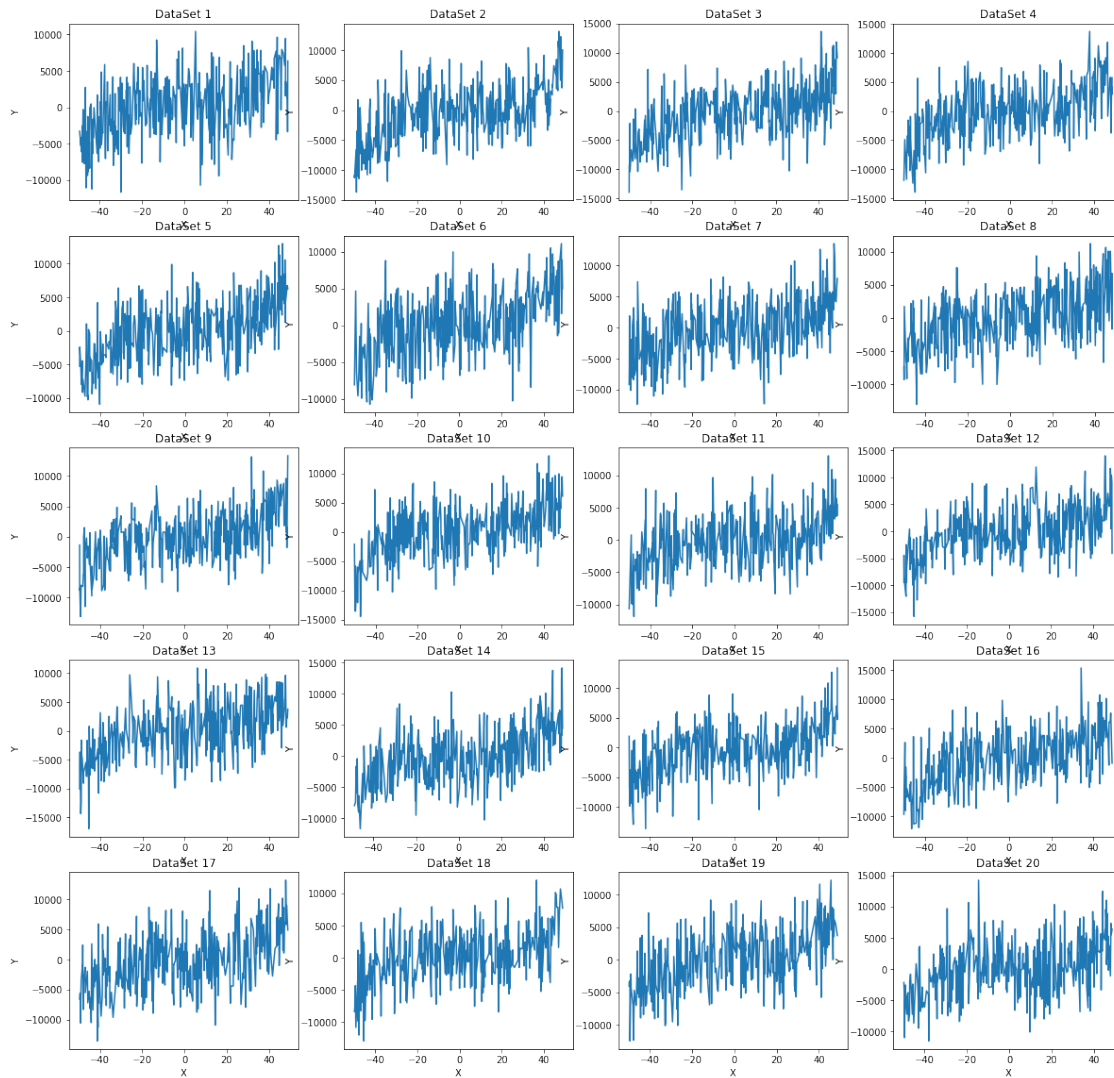
```
In [1]: import pickle
        from sklearn import model_selection, linear_model
        import numpy as np
        import pandas as pd
        import seaborn as sns
        from matplotlib import pyplot as plt

In [2]: xtrain = pickle.load(open('Q2_data/X_train.pkl', 'rb'))
        xtest = pickle.load(open('Q2_data/X_test.pkl', 'rb'))
        ytrain = pickle.load(open('Q2_data/Y_train.pkl', 'rb'))
        ytest = pickle.load(open('Q2_data/Fx_test.pkl', 'rb'))
        print('Shapes:', xtrain.shape, ytrain.shape, xtest.shape, ytest.shape)
```


Shapes: (20, 400) (20, 400) (80,) (80,)

The Subset division is done for us here. X and Y in the training set are already divided up into 20 groups with 400 examples each. Now let's plot some of those examples out.

```
In [3]: fig, ax = plt.subplots(5, 4, figsize=(20,20))
        ax = ax.reshape(-1)
        for i in range(20):
            sns.lineplot(xtrain[i], ytrain[i], ax=ax[i])
            ax[i].set_title('DataSet ' + str(i + 1))
            ax[i].set_xlabel('X')
            ax[i].set_ylabel('Y')
```



```
In [4]: sns.lineplot(xtest, ytest, color='Orange')
plt.title('Test Set')
plt.xlabel('X')
plt.ylabel('Y')
plt.show()
```



The Test Data is remarkably simple to model, yet the training data is just full of noise. I believe that there is underlying signal that the model can learn, though it's not clear to me. Most **higher order models will overfit this data**, that is evident. Every Model beyond Order 3 will have high variance.

```
In [5]: def k_poly(x: np.ndarray, k:int):
        return np.array([x ** i for i in range(k)]).T

In [6]: model = [
        [linear_model.LinearRegression().fit(
            k_poly(xtrain[i], order),
            ytrain[i])
         for i in range(20)
        ] for order in range(1, 10)
        ]

In [7]: # Use model[order - 1][train_group]
        def predict(x, group, order):
            assert 1 <= order <= 9 and 0 <= group < 20
            return model[order - 1][group].predict(k_poly(x, order))

        predict(np.array([10, 20, 30, 40, 50, 60, 70, 80, 90, 100]), group=2, order=2)

Out[7]: array([ 597.19529672, 1462.14120011, 2327.0871035 , 3192.03300688,
                4056.97891027, 4921.92481366, 5786.87071704, 6651.81662043,
                7516.76252382, 8381.7084272 ])
```

```

In [8]: import pandas as pd
        data = {'x_test': xtest, 'y_test': ytest}
        data.update({'g' + str(g) + '_o' + str(o): predict(xtest, g, o)
                      for o in range(1, 10) for g in range(20)})
        res = pd.DataFrame(data)
        res.head()

Out[8]:
```

	x_test	y_test	g0_o1	g1_o1	g2_o1	g3_o1	\
0	-21.47	-502.846044	172.115934	-295.486212	54.898088	-195.298503	
1	-37.11	-2581.477664	172.115934	-295.486212	54.898088	-195.298503	
2	-6.40	-12.862400	172.115934	-295.486212	54.898088	-195.298503	
3	3.61	3.055552	172.115934	-295.486212	54.898088	-195.298503	
4	25.77	843.144144	172.115934	-295.486212	54.898088	-195.298503	

	g4_o1	g5_o1	g6_o1	g7_o1	...	g10_o9	\
0	130.823889	221.373367	-285.143315	74.509785	...	-333.389085	
1	130.823889	221.373367	-285.143315	74.509785	...	-1641.092299	
2	130.823889	221.373367	-285.143315	74.509785	...	238.143596	
3	130.823889	221.373367	-285.143315	74.509785	...	256.042934	
4	130.823889	221.373367	-285.143315	74.509785	...	276.655114	

	g11_o9	g12_o9	g13_o9	g14_o9	g15_o9	\
0	56.274412	-218.516301	-818.700074	-296.567996	-190.478802	
1	-2887.226049	-2969.368153	-2575.118622	-2993.146071	-3794.663104	
2	-241.705202	-260.559395	-744.796127	-92.355826	144.453771	
3	3.458716	84.929019	-740.143493	-326.755585	571.694127	
4	846.448314	1116.700914	1027.585054	292.597502	999.181765	

	g16_o9	g17_o9	g18_o9	g19_o9
0	-413.194777	-223.125904	-703.640678	-289.987148
1	-3235.604606	-1425.033189	-2887.281348	-2722.331260
2	-130.127496	719.458849	901.813851	976.640329
3	-648.350204	416.465908	726.411816	99.572471
4	1054.439779	851.235641	634.307952	452.190180

[5 rows x 182 columns]

```

In [9]: def get_bias_variance(order):
        keys = ['g' + str(group) + '_o' + str(order) for group in range(20)]
        data = np.array([res[key].values for key in keys])
        variance = np.var(data, axis=0)
        bias = np.abs(np.mean(data, axis=0) - res['y_test'].values)
        return (np.mean(bias), np.mean(variance), np.sqrt(np.mean(bias ** 2)))

        get_bias_variance(order=1)

Out[9]: (1716.3588857838863, 20935.166423224564, 2600.4458653365477)

```



1.2 Analyzing Bias and Variance

We are plotting Bias and Variance for each model and against model size to see if the trends are satisfied. We expect that: * The Smaller Models will be High Bias and Low Variance since it has very little space to produce varying outputs, i.e. The function it learns will be simple. But bias is huge cause it didn't really learn a lot. * The Bigger Models will be High Variance and Low Bias since it can overfit the data getting the mean almost perfectly equal, but have huge deviations due to learning too complex a function on different inputs.


```
In [10]: bias_var_map = np.array([get_bias_variance(order)
                                   for order in range(1, 10)])
        bias_map = np.array([x[0] for x in bias_var_map])
        var_map = np.array([x[1] for x in bias_var_map])
        bssq_map = np.array([x[2] for x in bias_var_map])
        print('Bias:', bias_map, '\nVariance', var_map, '\nBias^2', bssq_map)
```

```
Bias: [1716.35888578  819.83786046  810.84021384   67.63398967   84.00189151
        79.18712742   80.09962645   84.97053521   85.57712158]
Variance [ 20935.16642322  70545.48914575 125870.85554877 150073.73954647
        212235.70832552 276388.48029738 316863.50071853 357510.87596821
        404293.66486518]
Bias^2 [2600.44586534  999.61412399  977.04619839   96.9006198   104.43825034
        96.63950685  101.23529961  101.66258248  100.73442981]
```

```
In [11]: sns.lineplot(np.arange(1, 10), bias_map, color='RoyalBlue')
        plt.ylabel('Bias')
        plt.xlabel('Degree')
        plt.title('Bias vs. Model Size')
        plt.show()
        sns.lineplot(np.arange(1, 10), var_map, color='IndianRed')
        plt.ylabel('Variance')
        plt.xlabel('Degree')
        plt.title('Variance vs. Model Size')
        plt.show()
```



```
In [12]: sns.lineplot(np.arange(1, 10), bssq_map + var_map, color='IndianRed')
sns.lineplot(np.arange(1, 10), bssq_map, color='RoyalBlue')
sns.lineplot(np.arange(1, 10), var_map, color='IndianRed')
plt.legend(['Total', '$Bias^2$', 'Variance'])
plt.ylabel('Error')
plt.xlabel('Degree')
plt.title('Bias vs. Variance')
plt.show()
```



1.3 Check the Model's Predictions

One final run to see how the models, on average are fitting the data

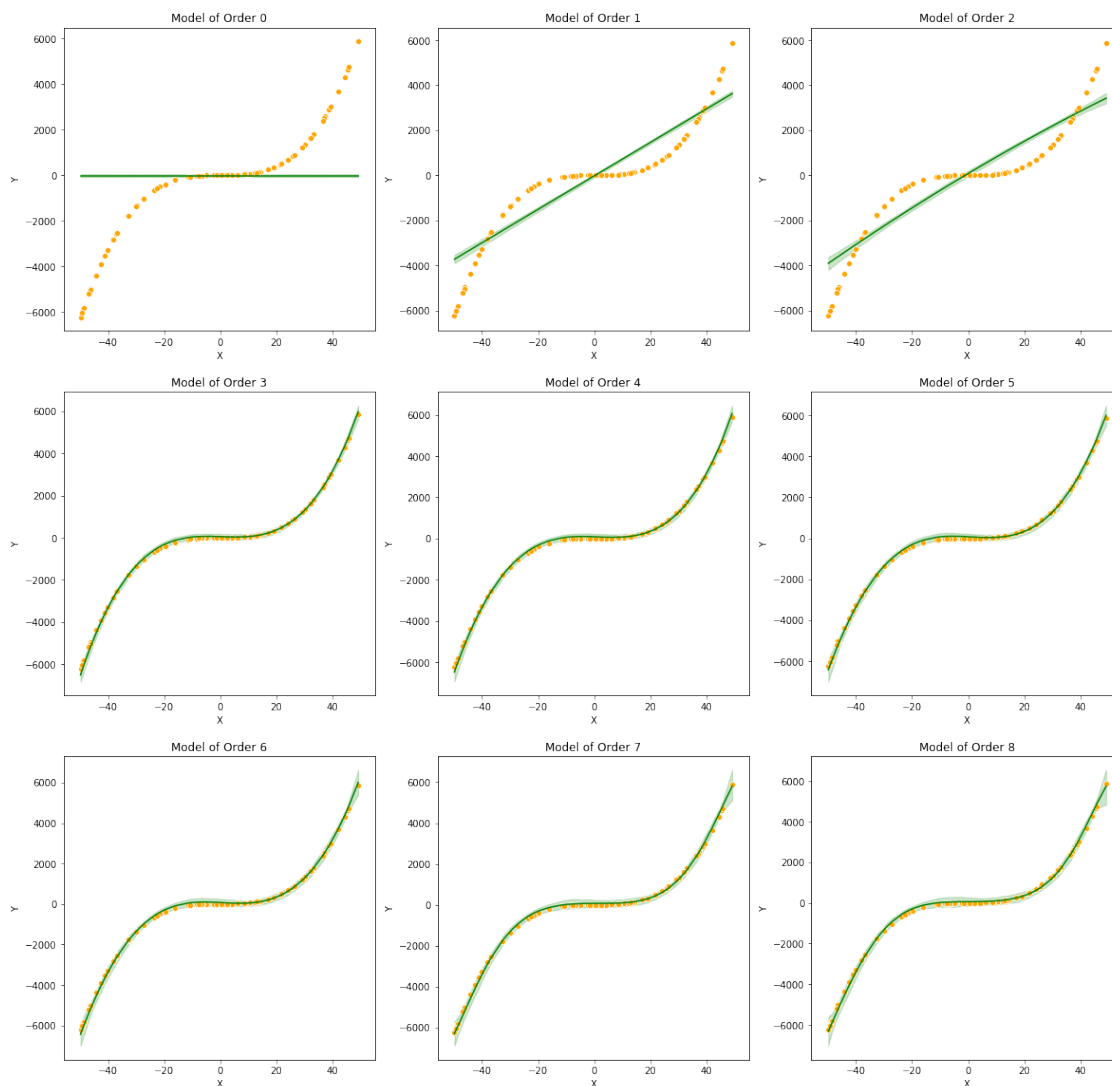
```
In [13]: fig, ax = plt.subplots(3, 3, figsize=(20,20))
ax = ax.reshape(-1)
print('Making Graphs: ', end='')
for order in range(1, 10):
    pred = np.array([model[order - 1][group].predict(k_poly(xtest, order))
                     for group in range(20)])
    ptest = np.mean(pred, axis=0)
    xfull, yfull = np.concatenate([xtest for i in range(20)]), pred.reshape(-1)
    print('.', end='')
    sns.lineplot(xfull, yfull, ax=ax[order-1], color='Green')
    sns.scatterplot(xtest, ytest, ax=ax[order-1], color='Orange')
```

```

ax[order - 1].set_title('Model of Order ' + str(order - 1))
ax[order - 1].set_xlabel('X')
ax[order - 1].set_ylabel('Y')
plt.show()

```

Making Graphs: ...



1.4 The Final Answers

Here the test set and the underlying structure of the data is of order 3. So our loss decreases till order 3 (actually till 4), then increases. Variance dominates later, Bias dominates first. So higher order models are (over 4) overfit and the lower orders models (below 3) are underfit.

Since the training set is very messy, we see higher order models gain variance pretty quickly. All data seems to be noise on the cubic form $y = x^3/32$.